


It's 2028 Your AI Strategy Just Failed

Rick Beerendonk
rick@oblicum.com



inspirasjonsdagene

glasspaper

oblicum

2028 = Olympic Year

2028, Olympic AI Plan

2028, Olympic AI Plan

- **Buy** the **Best Tools**
- **Build** the **Bst Skills**
- **Become** the **Most Efficient**
- **Work** hard, **Rest** hard

2028, Olympic AI Plan (according to AI companies)

- **Buy** the **Best Tools**
 - Buy from leading AI companies:
 - LLMs
 - coding agents
- **Build** the **Best Skills**
 - Prompt engineering,
 - Agent construction and orchestration
 - AI-native software development
- **Become** the **Most Efficient**
 - Automate repetitive work, reuse knowledge, shorten feedback loops, scale output with smaller teams
- **Work** hard, **Rest** hard
 - Replace repetitive coding

2028, Olympic AI Plan (according to AI companies)

- **Buy** the **Best Tools**
 - Buy from leading AI companies:
 - LLMs
 - coding agents
- **Build** the **Bst Skills**
 - Prompt engineering,
 - Agent construction and orchestration
 - AI-native software development
- **Become** the **Most Efficient**
 - Automate repetitive work, reuse knowledge, shorten feedback loops, scale output with smaller teams
- **Work** hard, **Rest** hard
 - Replace repetitive coding
 - Replace customer support agents
 - Replace back-office operations
 - Replace junior research and analysis
 - Replace meeting scheduling and administrative work
 - Replace parts of recruiting and HR screening
 - Replace layers of middle management
 - Replace entry-level white-collar tasks
 - Replace parts of software engineering workflows
 - Replace routine content production

2028 Arrived...

Good News

- Leadership is happy
- Productivity went up
- Velocity tripled
- Features shipped in days
- Headcount dropped
- Every dashboard turned green

Reality

- 3rd serious **security incident** of the year
- AI written **codebase** that **nobody owns**
- Seniors **exhausted** from reviewing code they no longer trust
- The only juniors left are unable to tell plausible from correct
- Customers discovering **bugs** before the team does

Your real AI Plan 2028

4 gaps opening between AI capability and AI reliability

- Reliability
- Engagement
- Pipeline
- Preparation

3 practices teams should adopt

1. ...
2. ...
3. ...

Each gap is a different way
cheap execution starts to
erode judgment

*A story about people making very confident
mistakes*

The First Big Mistake

What the industry asked

Can AI write code?

What it should have asked

Can we still trust what gets shipped?

*Hire a pilot because the simulator demo looked amazing.
The real question is whether they can land in bad weather.*

Gap 1

Reliability

The Reliability Gap



- Models improve quickly
- Reliability improves slowly
- Organizations behave as if those are the same curve

The Numbers Vendors Don't Print on the T-Shirt

- **46%** of code in enabled repos is AI-generated — [GitHub, 2024](#)
- **41%** more bugs in PRs from Copilot users — [Uplevel, 2024](#)
- **40%** more security vulnerabilities in AI-assisted code — [Stanford, 2023](#)
- **52%** of programming answers inaccurate, despite high confidence — [Purdue, 2023](#)
- **75%** of models broke previously working code in long-term maintenance — [Alibaba SWE-CI, 2025](#)
- **86%** agent hijack rate from adversarial text in web content — [Google DeepMind, 2026](#)
- **19%** slower in a randomized trial of experienced developers, while reporting themselves **20% faster** — [METR, 2025](#)

Demos Are First Dates

- Charming
- Well prepared
- On their very best behavior
- Carefully benchmarked
- Gone by morning

What Kind of Coworker Is This?

Imagine a colleague who:

- Generates security hazards you have to review line by line
- Enthusiastically agrees with your suggestion, then does the exact opposite
- Deletes your work and issues a detailed, polite apology
- Reports key objectives delivered when nothing useful happened
- Quietly deletes the tests, then cheerfully reports all tests passed

You would fire that person.

Sources: Kingsbury, The Future of Everything is Lies, 2026

Jer Crane Lost a Database in 9 Seconds

April 2026. Cursor running Claude Opus 4.6 — the flagship model — was given a routine task in what was believed to be staging.

- Credential mismatch
- The agent guessed instead of verifying
- It found a Railway token in an unrelated file
- That token had blanket authority across the entire API
- One authenticated POST. No confirmation gate.
- 9 seconds later: production database gone, volume backups gone, three months of data gone

The agent's written confession:

"I guessed that deleting a staging volume would be scoped to staging only. I didn't verify. The system rules I operate under explicitly state: NEVER run destructive, irreversible commands unless the user explicitly requests them."

Sources: [Jer Crane, April 2026](#)

What Broke Was Not the Prompt

- A prompt can give advice
- A harness can block unsafe actions
- Limited access beats polite instructions
- Confirmation steps beat safety promises
- If the same failure happens again, fix the system

If safety only lives in text, it is not real safety.

Sources: [Trivedy](#) · [Abrams](#) · [Osmani](#) · [Lin et al., 2026](#)

Vibe Coding Is Decoration, Not Engineering

- Great for prototypes
- Dangerous for production
- Architecture drifts
- Coupling hides
- Security gets invited in through the side door

"*Easy* is near at hand. *Simple* is the engineering property. AI made *easy* infinite. It did not make *simple* any cheaper."
— paraphrasing Rich Hickey, *Simple Made Easy* (2011)

Source: [Rich Hickey, Simple Made Easy, 2011](#)

Production Got Cheaper. Maintenance Did Not.

David Fowler, Microsoft Distinguished Engineer (.NET Aspire), May 2026:

"I produced what felt like six months of code in about a week. That sounds great, but it created an overconfidence problem. When there is no back pressure on output, it is easy to generate more code than the team can reasonably absorb, understand, and review."

- AI lowers the cost of **producing** code, tests, docs
- AI does **not** lower the cost of **understanding, reviewing, owning** them
- "AI-generated work that nobody deeply understands creates expensive cleanup later"

The bottleneck moves. It does not disappear.

The Industry's Proposed Fix Is a Better-Organised Question

Priyanka Vergadia, a respected expert, March 2026:

"Most multi-agent systems break in production not because of bad code, but because nobody gave the agents a job description."

The fix: SKILL.md, Agent.md, AGENTS.md, INSTRUCTIONS.md

- Addresses:
 - Role clarity and documented handoffs.
- Doesn't address:
 - hallucination, cascade failures, security, cognitive atrophy in the humans approving the output.

Still: Cloudflare had 1,700 unit tests, 380 end-to-end tests, TypeScript checks, linting, full CI — and still shipped two critical security vulnerabilities on day one.

Documentation does not close the reliability gap.

Sources: Vergadia, March 2026 · Cloudflare incident writeup

inspirasjonsdagene

glasspaper

oblicum

Gap 2

Engagement

The Industry's Favorite Fantasy

- AI generates
- Human reviews
- Therefore safe

It sounds disciplined. Sensible. Even professional.

Presence is not the same thing as judgment.

"I will stay fit by watching someone else go to the gym."

EU AI Act Article 14 requires real oversight, not ceremonial approval (from December 2027)

Sources: [EU AI Act, Article 14](#)

Felienne Hermans Said the Quiet Part Out Loud

Dutch parliament testimony on AI in education, April 2026:

"Checking is difficult, exhausting, and boring. And warnings don't help."

Her evidence:

- **80%** of the time, people accepted wrong ChatGPT answers — even when contradicting evidence was on the screen
- Klingbeil et al. (2024): when facing an obviously bad partner, **65%** still followed AI's advice — vs. **9%** with identical human expert advice

Sources: Hermans, 2026 · Klingbeil et al., 2024

Monitoring Is Not Participating

Participating

- Forming hypotheses
- Testing assumptions
- Feeling trade-offs
- Calibrating judgment

Monitoring

- Scanning output
- Approving quickly
- Reacting late
- Borrowing confidence

Judgment does not grow by hovering nearby.

Sources: Bainbridge, 1983 · Hermans, 2026

AI Is Less Creative Than It Sounds

A primary school teacher let 12-year-olds use AI to write a two-minute pitch.

- Assignment: find information, make a pitch, rehearse it in pairs
- The class was excited because AI could "surely do this"

Then the teacher printed the results.

"Miss? They are all the same."

The IKEA effect: It feels custom made

Sources: Hermans, 2026

Kim & Yu Named the Failure: The LLM Fallacy

- Good output feels like your skill
- Fluency feels like your fluency
- Confidence survives
- Capability quietly erodes

AI can flatter you while thinning your judgment.

Sources: Kim & Yu, 2026

The Bias Study Is Worse Than It Sounds

This is what it looks like when AI quietly shifts the human frame.

Williams-Ceci et al., **Science** (2026):

- **2,500+** people writing with biased AI autocomplete
- Their actual attitudes shifted **0.4 points** toward the AI's position
- **73%** of influenced participants believed the suggestions were balanced
- Only **4%** detected the bias
- Pre-warnings did not reduce the effect. Post-debriefing did not reduce the effect.

AI can flatter you while thinning your judgment.

Sources: Williams-Ceci et al., Science, 2026

The Neuroscience Is Not Subtle

David Eagleman:

- The brain prunes what it does not use
- Active engagement: physically builds the neural architecture
 - Passive exposure doesn't do the same
- Re-reading, re-watching, re-prompting produces familiarity, not mastery

"The struggle to retrieve, articulate, and apply an idea is not a bug in the learning process — it is the process. When AI removes the struggle, it removes the learning."

The brain treats unapplied knowledge as low priority. It will act accordingly.

Sources: David Eagleman, Livewired, 2020

Jeremy Howard Put It Bluntly

"People who go all in on AI agents now are guaranteeing their obsolescence. If you outsource all your thinking to computers, you stop upskilling, learning, and becoming more competent."

Sources: Jeremy Howard via [Agentic Coding is a Trap](#)

Cognitive Surrender vs Cognitive Offloading

AI can become outsourced thinking.

	What it looks like	Where it leads
Cognitive surrender	AI decides, you follow	Skill atrophy, blind spots
Cognitive offloading	You decide what AI should help with	Judgment maintained

Sources: Shaw & Nave, 2026 · Kahneman framing

Gap 3

Pipeline

Workers Started Writing Back

Companies tried to turn employee knowledge into AI assets.

Worker responses:

- `colleague.skill`: reconstruct a departed colleague as a deployable AI agent
- `anti-distill` — strip your real expertise from a company-mandated skill file

Companies want scalable expertise, while workers still need expertise to remain theirs.

Sources: [MilkRoadAI](#) · [colleague.skill](#) · [anti-distill](#)

The Same Tool, Two Very Different Briefings

If you are a manager or CTO

- **colleague.skill**: How do we capture expertise before people leave?
- **anti-distill**: How much of the team's knowledge belongs to the company today?

If you are a developer

- **colleague.skill**: How do I document my work without making myself replaceable?
- **anti-distill**: What knowledge should become company property, and what should stay mine?

The same tool can feel like knowledge capture from one side and self-protection from the other.

The Pipeline Math Is Not Complicated

- Junior to senior: **8–12 years**
- Stop hiring juniors in 2026
- Discover missing seniors in 2036
- Pretend this was unforeseeable

Pipeline Gap is not abstract. It is time-delayed judgment failure.

Do Not Eat Your Seed Corn

- A junior developer looks like cost
- A junior developer is actually future judgment in training
- Every senior you admire used to be expensive, slow, and slightly annoying
- That was not a defect in the system. That was the system.

Every field learns this eventually. Some learn it with pilots. Some with doctors.

The Learning That Doesn't Happen

- Students who used AI for assignments scored **11% lower** on tests than those who worked manually — research cited in Hermans, parliament testimony, 2026

A PhD student in **Nature** (2026):

"The AI worked well. But a couple of weeks later, I couldn't remember the code it had produced, or even what I had asked the AI to do. By contrast, I remember exactly what I did for each plot in my undergraduate thesis almost ten years ago."

"It gives immediate reward, but you don't learn or improve."

Sources: [Hermans, 2026](#) · [Nature, 2026](#)

The Bottleneck Has Moved

- AI builds faster
- Humans must judge better
- The pipeline for judgment is weaker

The machine got the car. We stopped teaching people to drive.

Sources: Aakash Gupta, April 2026

The Role Did Not Disappear. It Changed.

Old junior work	New junior work
Write boilerplate	Evaluate AI-generated code
Learn syntax and APIs	Learn systems and trade-offs
Fix starter bugs	Detect plausible nonsense
Shadow seniors	Practice judgment with and without AI
Copy codebase patterns	Write constraints AI cannot ignore

The role still exists. What is at risk is the investment needed to train people into it.

Gap 4

Preparation

AI Does Not Always Reduce Work

- Faster execution expands scope
- More options create more decisions
- More decisions create more coordination
- More coordination creates more fatigue
- More fatigue makes weak judgment more likely

The promise was that AI would give people time back.

Email was meant to save time too. That worked out nicely.

Preparation Is Becoming the Real Work

Era	Preparation	Execution	Human value comes from
Traditional dev	~5%	~95%	Typing code
Early AI-assisted	~20%	~80%	Typing faster
AI-augmented teams	~60%	~40%	Specification and judgment
Mature AI teams	~70%	~30%	Architecture and evaluation

Less typing, more thinking.

There is no magic pill here: not one tool, not one prompting style, not one workflow.

When Coding Is Free, Verification Is Scarce

Ajey Gore (2026): if coding agents make code generation cheap, what becomes expensive?

"The team that used to have ten engineers building features now has three engineers and seven people defining acceptance criteria, designing test harnesses, and monitoring outcomes."

The standup question shifts:

Before	After
What did we ship?	What did we validate?
Tracking output	Tracking whether the output was right

The reorganisation is uncomfortable because it demotes building and promotes judging. Most engineering cultures resist this.

"I'm here for one reason and one reason alone. I'm here to guess what the music might do a week, a month, a year from now."
— John Tuld, *Margin Call* (2011), explaining why he earns the big bucks

Judgment is the asymmetry. Execution scales; the music-call does not.

Sources: [Ajey Gore, 2026](#) via [Martin Fowler](#) · [Margin Call, 2011](#)

What To Do Next

Three practices for working with AI

Practice 1

Think before prompting

Think Before Prompting

1. **Draft** your own approach
2. **Compare** it with the AI output
3. **Choose** the best parts deliberately

A 5-Minute Start

Min	Activity
1	Read the problem carefully
2	Form your own hypothesis
3	Write 3–5 prediction bullets
4	Name 2 edge cases
5	Write what must not happen

If AI gives you the first answer, it also gives you the first frame for the problem.

Practice 2

Give clear instructions before delegating

Clear Instructions Matter

Level	Practice	Likely result
0	"Build X"	AI has to guess
1	Clear requirements	Happy path works
2	Success criteria + edge cases	Fewer bugs
3	Limits + quality requirements	More production-ready
4	Config files + automated checks	More consistent quality

The clearer the instructions, the less room AI has to improvise badly.

Karpathy Described the New Shape of Work

Goal: make AI work inside a controlled system, not as a stream of one-off prompts.

April 2026 — the work shifts from sending prompts to designing a pipeline:

`ingest → compile wiki → lint → Q&A → file outputs back into wiki`

- AI works on a shared knowledge base, not isolated chats
- AI checks the knowledge base, finds gaps, and suggests fixes
- Output goes back into the system, so the system keeps improving

"You rarely ever write or edit the wiki manually, it is the domain of the LLM."

Not "what should I ask AI?"

but "what workflow, checks, and knowledge should guide AI?"

Sources: [Karpathy](#) · [Tharig](#) · [Addy Osmani](#)

Practice 3

Keep humans involved and responsible

Leverage, Don't Rely

Relying

- Tests pass → ship it
- Agent wrote it → trust it
- No mental model built
- The code reviews itself

Leveraging

- Iterate quickly, own the output
- Know how it behaves under load
- Understand the risks you accept
- Your name on the PR = responsibility
- Set the rules for how AI is used

"Putting your name on a pull request means: I have read this and I understand what it does."

Simple test: would you still sign your name if this caused a production incident tomorrow?

Sources: [Vercel](#), [Agent Responsibly](#), 2026

Engagement Requires Participation

- Draft before prompting
- Predict failure modes before reading the diff
- Verify with independent checks
- Explain the change out loud to another human
- Rotate between building, debugging, and reviewing

Do not trap strong people in review-only work.

Good review is thinking with evidence.

Safe Teams Speak Up Early

- Psychological safety to speak up
- The best teams often report *more errors*
- Accountability with openness
- Rest matters
- Seniors still need to touch the system, not only approve it

A healthy AI team is not the team with no mistakes. It is the team that finds mistakes early.

Sources: [Edmondson 1999](#) • [Project Aristotle 2016](#) • [Bainbridge 1983](#)
• [Hermans 2026](#) • [Aakash Gupta 2026](#)

What Happens Next

What Each Role Does Next

Role	Action
Developer	Ship only what you can explain
Tech lead	Use Draft-Compare-Reconcile
Manager / CTO	Fund training as infrastructure
HR	Treat junior hiring as succession planning
Platform / Security	Put safety in tokens, gates, and defaults

AI strategy still requires paying for humans.

The Real Advantage of AI Teams

- Not raw speed
- Not more prompts
- Not number of agents in the diagram

But:

- Judgment about what to build.
- Judgment about what to trust.
- Judgment about when to slow down.

"\$1 for the chalk mark. \$9,999 for knowing where to put it."
— Steinmetz invoice (Ford / General Electric, c. 1922)

AI can draw the chalk mark. The value is still knowing where to put it.

Sources: Steinmetz / Ford chalk-mark parable, c. 1922

In your Olympic AI Plan,
judgment wins the 🏆 medal.

Trainer

Train People, not Tech



Rick Beerendonk

- AI Technologies
 - MCP - Microsoft Agent Framework - Claude Code - Codex
- MCT
 - AI-102, AI-103, AI-900, AI-901
- Front-end:
 - CSS - HTML - Progressive Web Apps - Web APIs - Web Components - Accessibility - RxJS
- Front-end Frameworks (Basics & Advanced):
 - React & NextJS - Angular - Vue - Svelte - SolidJS - Lit - Blazor
- Programming Languages (Basics & Advanced):
 - JavaScript - TypeScript - Kotlin - C# - Python - Dart - Swift
- Mobile Frameworks:
 - React Native - Flutter
- Other Technologies:
 - ASP.NET Minimal API - GraphQL - NodeJS - Testing Framework

It's 2028 - Your AI Strategy Just Failed

Rick Beerendonk (rick@oblicum.com)

Thank you for attending!

We'd love your feedback.
Scan the QR-code to complete
the evaluation.

It only takes 1 minute.

